Watsonx.data

Day 1
Fundamentals

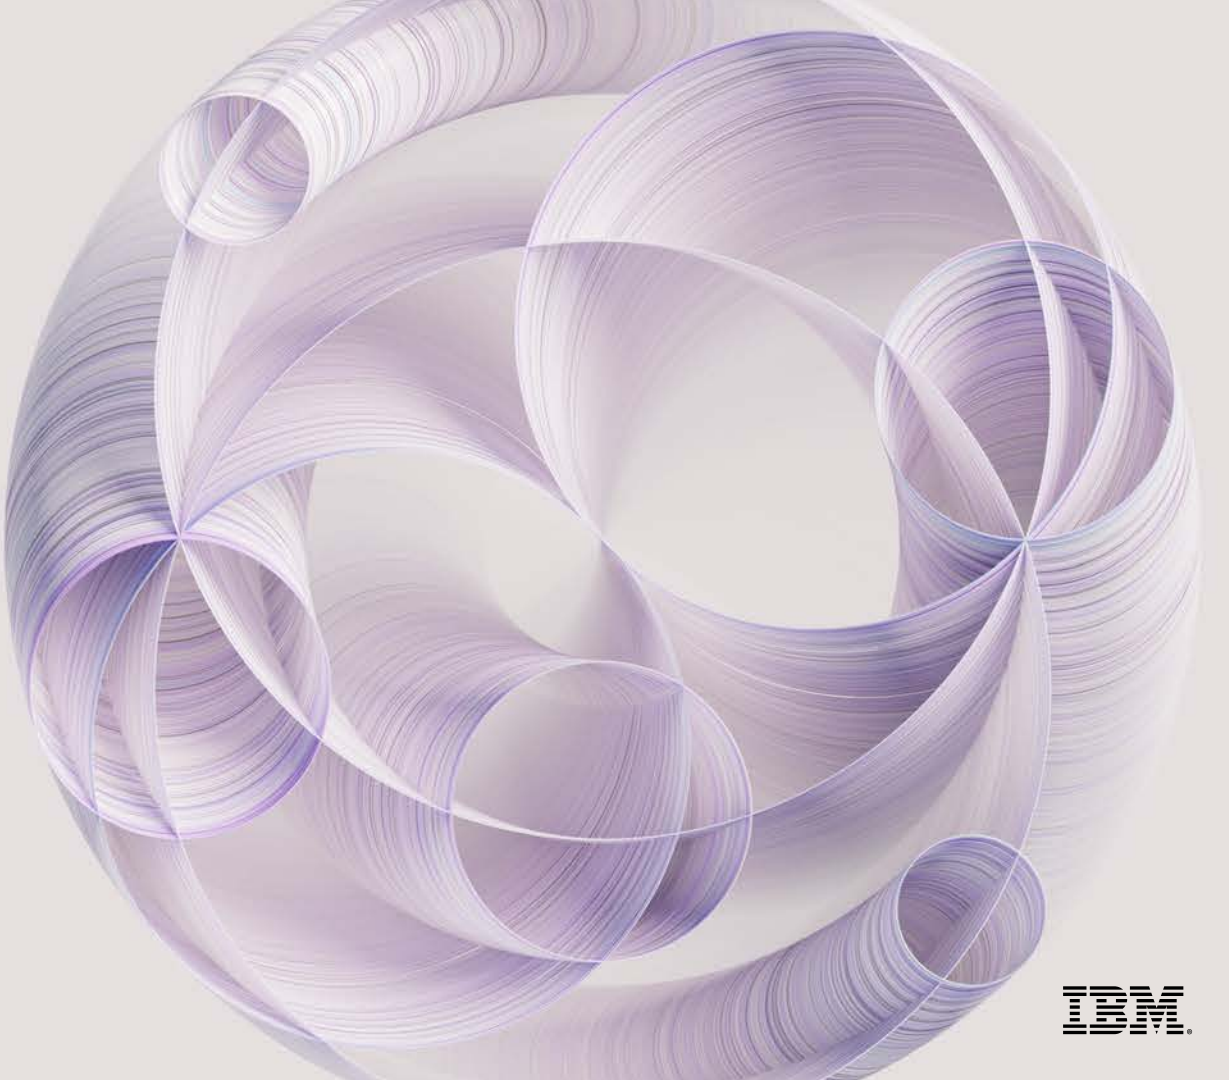

Access all your data across hybrid-cloud through a single point of entry

An open data store built for hybrid deployment of your analytics and AI workloads

1. Share a single copy of data with tools that can read open data formats to minimize data duplication

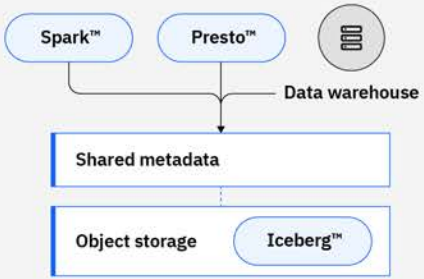

2. Connect to and access data remotely across hybrid-cloud with the ability to cache remote sources

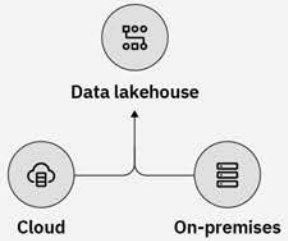

3. Synchronize and incorporate Db2 for z/OS data for lakehouse analytics.

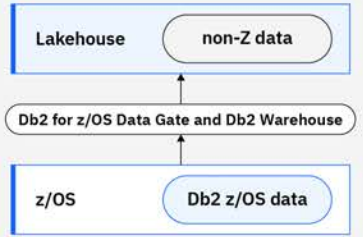

Get started in minutes with built-in governance, security and automation.

Accelerate time to trusted analytics and AI

Connect to your existing analytics data and deploy fit-for-purpose engines in minutes

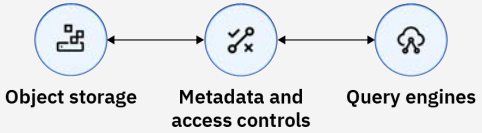

Address enterprise compliance and security using built-in centralized governance across your data ecosystem

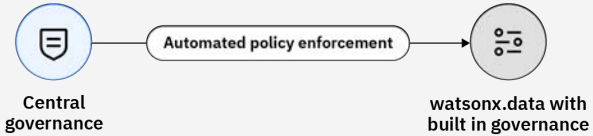

Use foundation models to discover, augment, refine, and visualize watsonx.data data and metadata

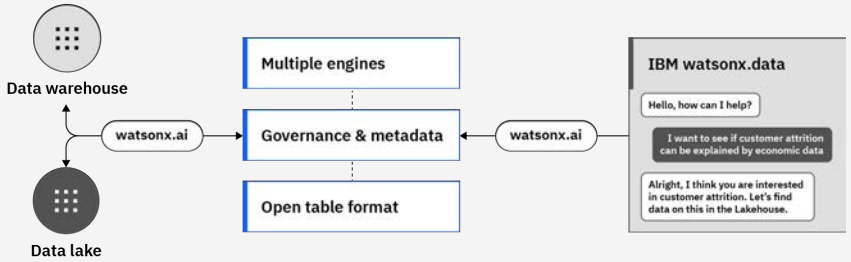

Reduce your data warehouse costs by up to 50%* by optimizing workloads

Optimize workloads from your data warehouse when you take advantage of low-cost object storage and fit-for-purpose query engines

*When comparing published 2023 list prices normalized for VPC hours of IBM watsonx.data to several major cloud data warehouse vendors. Savings may vary depending on configurations, workloads and vendors.

1 Share data between multiple analytics engines

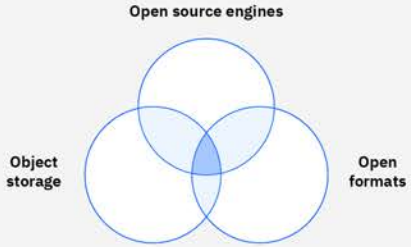

2 Use fit-for-purpose compute and cache-optimized instances

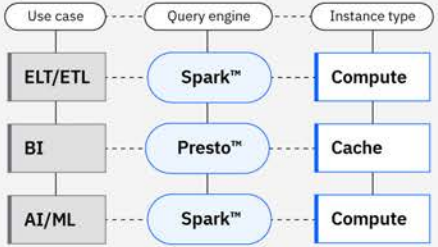

3 Scale up and scale down automatically

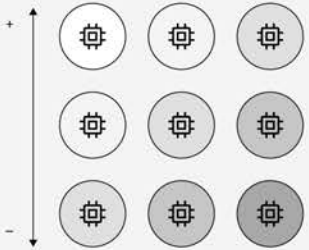

Common data file formats

Computer systems and applications store data in files

Data can be stored in binary or text format

File formats can be open or closed (proprietary/lock-in)

Open formats (Parquet, ORC, and Avro) are commonly used in data lakes and lakehouses

CSV

- Human-readable text
- Each row corresponds to a single data record
- Each record consists of one or more fields, delimited by commas

{ JSON }

- Human-readable text
- Open file and data interchange format
- Consists of attribute-value pairs and arrays
- JSON = JavaScript Object Notation

Parquet

- Open-source
- Binary columnar storage
- Designed for efficient data storage and fast retrieval
- Highly compressible
- Self-describing

Apache orc

- Open-source
- Binary columnar storage
- Designed and optimized for Hive data
- Self-describing
- Similar in concept to Parquet

AVRO

- Open-source
- Row-oriented data format and serialization framework
- Robust support for schema evolution
- Mix of text/binary

THE OPEN DATA LAKEHOUSE DIFFERENTIATOR

“IBM / Cloudera Shares a Joint Vision”

The “Old” Way

Hive SQL over data in HDFS

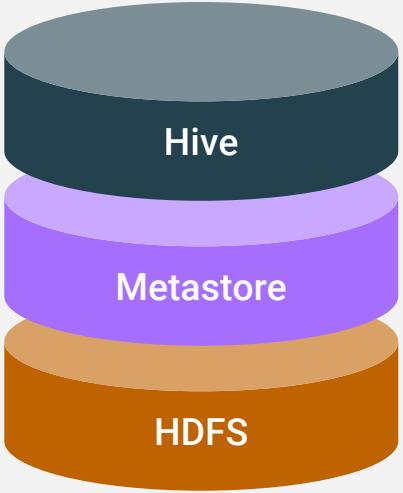

“Some” other way

Spark SQL over Delta tables

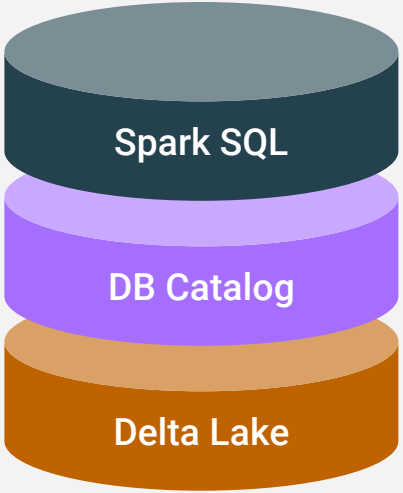

The Iceberg way

Multi-function analytics over all your data

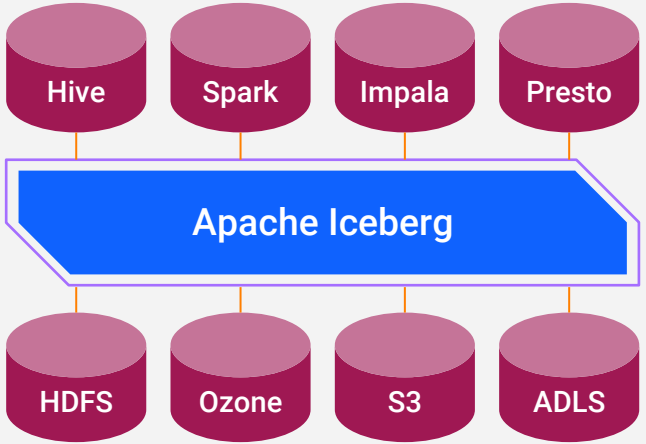

Powerful Engines

Open-source Engines for each use case

presto Apache Spark

04

Catalog

IBM use Hive Metastore to catalog Iceberg Table Formats

03

Table Format

Organize the data where it lives into tables using Iceberg

02

File Format

Store Data in Apache Parquet Files (Open Columnar Format)

01

Storage

amazon S3 IBM Cloud Object Storage

Hyperscaler Provider (AWS, Azure or IBM)

Which cluster type?

Depending on your use-case!

There are three models you should consider for your cluster:

- Compute
- Balanced
- Storage

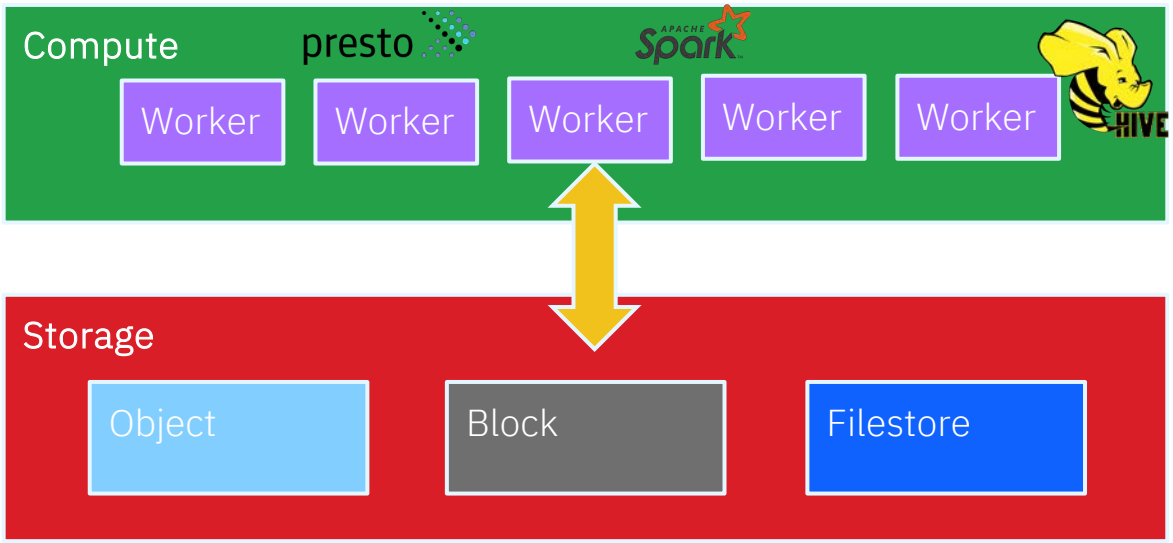

Which storage?

Recommendations for Storage Decision